

# The Chow-Liu Algorithm

C. K. Chow and C. N. Liu. Approximating discrete probability distributions with dependence trees. IEEE Transactions of Information Theory, IT-14(3), 1968.

1

## The Goal

Given a finite set of samples in a dataset, estimate the underlying  $n$ -dimensional discrete probability distribution using a **tree model**.

2

## Trees

What is a tree?


- The variables in the dataset are the vertices  $V$
- There are edges in the set  $E$  that connect the vertices
- We'll assume the edges are undirected for now
- A graph  $(V, E)$  is a tree if it is connected and has no cycles

Technical point: We will allow our trees to be a forest ie. the tree model we learn may be disconnected

3

## Trees

- In a directed tree, we pick a vertex as the root
- We then turn the edges into directed edges and orient the edges away from the root
- This means that each vertex has at most one parent (but may have more than one child)



4

## Tree Models

Notation:

- $\mathbf{x}$  (as in bold x) is an n-dimensional vector ie.  $\mathbf{x} = (x_1, x_2, \dots, x_n)$
- Each  $x_i$  in  $\mathbf{x}$  is a variable
- $P(\mathbf{x})$  is a joint probability distribution of  $n$  discrete variables  $x_1, x_2, \dots, x_n$

5

## Tree Models

- We want to approximate the true joint probability distribution using tree models of the form:

$$P_t(x) = \prod_{i=1}^n P(x_i | x_{\pi(i)})$$

- $\pi(i)$  means “parent of variable i”
- If  $i$  is the root then  $\pi(i)$  is the empty set:  
 $P(x_i | x_{\pi(i)}) = P(x_i)$

6

## Tree Models

$$P_t(x) = \prod_{i=1}^n \underbrace{P(x_i | x_{\pi(i)})}_{\text{pairwise}}$$

- Tree models consider the **pairwise** relationships between variables in the dataset
- It is an improvement over just treating the variables independently of each other

7

## Closeness of approximation

- Let  $P(\mathbf{x})$  and  $P_t(\mathbf{x})$  be two probability distributions of  $n$  discrete variables  $\mathbf{x} = (x_1, x_2, \dots, x_n)$ .
- Let

$$KL(P, P_t) = \sum_{\mathbf{x}} P(\mathbf{x}) \log \frac{P(\mathbf{x})}{P_t(\mathbf{x})}$$

Note: This summation is over all configurations of  $(x_1, x_2, \dots, x_n)$

The formula for  $KL(P, P_t)$  is called the Kullback-Leibler divergence (or KL divergence for short)

8

## Kullback-Leibler Divergence

- We'll rewrite the KL divergence as:

$$KL(P, P_t) = \sum_x P(x) \log P(x) - \sum_x P(x) \log P_t(x)$$

- The first term doesn't depend on  $P_t$ .
- The second term is known as the cross-entropy between  $P$  and  $P_t$ .
- Properties of KL divergence:
  - $KL(P, P_t) \geq 0$
  - $KL(P, P_t) = 0$  if and only if  $P(x) \equiv P_t(x)$  for all  $x$

9

## A Minimization Problem

Given:

- An  $n$ th-order probability distribution  $P(x_1, x_2, \dots, x_n)$  with  $x_i$  being discrete
- $T_n$ - The set of all possible first-order dependence trees

Find the optimal first-order dependence tree  $\tau$  such that  $KL(P, P_\tau) \leq KL(P, P_t)$  for all  $t \in T_n$ .

10

## Exhaustive Search

- Why not just search over all possible trees?
- Not feasible -- there are  $n^{(n-2)}$  possible trees with  $n$  vertices (from Cayley's formula)
- We will turn the search into a maximum weight spanning tree (MWST) problem

11

## Mutual Information

- Define the mutual information  $I(x_i, x_j)$  between two variables  $x_i$  and  $x_j$  to be:

$$I(x_i, x_j) = \sum_{x_i, x_j} P(x_i, x_j) \log \left( \frac{P(x_i, x_j)}{P(x_i)P(x_j)} \right)$$

- Key insight: a probability distribution of tree dependence  $P_t(\mathbf{x})$  is an optimum approximation to  $P(\mathbf{x})$  iff its tree model has maximum weight
- Proof to follow

12

## Proof

$$\begin{aligned}
 KL(P, P_t) &= \sum_{\mathbf{x}} P(\mathbf{x}) \log P(\mathbf{x}) - \sum_{\mathbf{x}} P(\mathbf{x}) \sum_{i=1}^n \log P(x_i | x_{\pi(i)}) \\
 &= \sum_{\mathbf{x}} P(\mathbf{x}) \log P(\mathbf{x}) - \sum_{\mathbf{x}} P(\mathbf{x}) \sum_{i=1, \neq \text{root}}^n \log \frac{P(x_i, x_{\pi(i)})}{P(x_{\pi(i)})} \\
 &= \sum_{\mathbf{x}} P(\mathbf{x}) \log P(\mathbf{x}) - \sum_{\mathbf{x}} P(\mathbf{x}) \sum_{i=1, \neq \text{root}}^n \log \frac{P(x_i, x_{\pi(i)})}{P(x_i) P(x_{\pi(i)})} \\
 &\quad - \sum_{\mathbf{x}} P(\mathbf{x}) \sum_{i=1}^n \log P(x_i)
 \end{aligned}$$

13

## Proof (continued)

Note that:  $-\sum_{\mathbf{x}} P(\mathbf{x}) \log P(x_i) = -\sum_{x_i} P(x_i) \log P(x_i)$

To see this, suppose  $\mathbf{x} = (x_1, x_2)$ , let all variables are binary, let  $i=1$

$$\begin{aligned}
 &-\sum_{\mathbf{x}} P(\mathbf{x}) \log P(x_i) \\
 &= -[P(x_1 = 0, x_2 = 0) \log P(x_1 = 0) + P(x_1 = 0, x_2 = 1) \log P(x_1 = 0) + \\
 &\quad P(x_1 = 1, x_2 = 0) \log P(x_1 = 1) + P(x_1 = 1, x_2 = 1) \log P(x_1 = 1)] \\
 &= -[P(x_1 = 0) \log P(x_1 = 0) + P(x_1 = 1) \log P(x_1 = 1)] \\
 &= -\sum_{x_1} P(x_1) \log P(x_1) = -\sum_{x_i} P(x_i) \log P(x_i)
 \end{aligned}$$

14

## Proof (continued)

In the same way:

$$\begin{aligned} & \sum_x P(x) \log \frac{P(x_i, x_{\pi(i)})}{P(x_i)P(x_{\pi(i)})} \\ &= \sum_{x_i, x_{\pi(i)}} P(x_i, x_{\pi(i)}) \log \frac{P(x_i, x_{\pi(i)})}{P(x_i)P(x_{\pi(i)})} = I(x_i, x_{\pi(i)}) \end{aligned}$$

15

## Proof (continued)

One more piece of notation:

$$\begin{aligned} H(\mathbf{x}) &= - \sum_x P(\mathbf{x}) \log P(\mathbf{x}) \\ H(x_i) &= - \sum_{x_i} P(x_i) \log P(x_i) \end{aligned}$$

Substituting the expressions above and from pg 12 into the last line of pg 13:

$$KL(P, P_t) = - \sum_{i=1}^n I(x_i, x_{\pi(i)}) + \sum_{i=1}^n H(x_i) - H(\mathbf{x})$$

16



## Proof

$$KL(P, P_t) = - \sum_{i=1}^n \underbrace{I(x_i, x_{\pi(i)})}_{\text{Mutual information is always } \geq 0} + \sum_{i=1}^n \underbrace{H(x_i) - H(\mathbf{x})}_{\text{Independent of the dependence tree}}$$

Mutual information  
is always  $\geq 0$

Independent of the  
dependence tree

Minimizing  $I(P, P_t)$  is the same as maximizing the total branch weight:

$$\sum_{i=1}^n I(x_i, x_{\pi(i)})$$

17

## The algorithm

- First calculate all  $n(n-1)/2$  pairwise mutual information measures
- Use Kruskal's algorithm to construct maximum weight spanning tree:
  - Construct tree one edge at a time, in decreasing order of the weights
  - If all weights are  $> 0$ , you get one connected component
  - Running time is  $O(n^2)$  for  $n$  variables because you have to consider all  $n(n-1)/2$  edges

18

## Estimation

- But in order to calculate mutual information  $I(x_i, x_j)$ , you need the probability distribution  $P(\mathbf{x})$
- Need to estimate the mutual information from a finite set of samples using maximum likelihood estimation

19

## Estimation

Suppose you are given  $s$  independent samples  $\mathbf{x}^1, \mathbf{x}^2, \dots, \mathbf{x}^s$  of a discrete variable  $\mathbf{x}$ . Each sample is an  $n$ -component vector ie.  $\mathbf{x}^k = (x_1^k, x_2^k, \dots, x_n^k)$ .

Define:

$n_{uv}(i, j) = \# \text{ of samples with } x_i = u \text{ and } x_j = v$

$$f_{uv}(i, j) = \frac{n_{uv}(i, j)}{\sum_{u,v} n_{uv}(i, j)}$$

Maximum Likelihood  
Estimator for  $P(x_i = u, x_j = v)$

$$f_u(i) = \sum_v f_{uv}(i, j)$$

Maximum Likelihood  
Estimator for  $P(x_i = u)$

20

## Estimation

Calculate:

$$\hat{I}(x_i, x_j) = \sum_{u,v} f_{uv}(i, j) \log \frac{f_{uv}(i, j)}{f_u(i) f_v(j)}$$

Use  $\hat{I}(x_i, x_j)$  in Kruskal's algorithm instead of  $I(x_i, x_j)$

21

## The entire algorithm

1. Compute marginal counts  $f_u(i)$  and pairwise counts  $f_{uv}(i, j)$
2. Compute mutual information  $\hat{I}(x_i, x_j)$  for all pairs  $x_i$  and  $x_j$
3. Compute MWST using Kruskal's algorithm. Pick a root, orient edges away from the root.
4. Set the parameters in the CPTs for each node to be their maximum likelihood estimates:

$$P(x_i | x_{\pi(i)}) = \frac{f_{uv}(i, \pi(j))}{f_u(i)}$$

22

## The entire algorithm

1. Compute marginal counts  $f_u(i)$  and pairwise counts  $f_{uv}(i,j)$
2. Compute mutual information  $\hat{I}(x_i, x_j)$  for all pairs  $x_i$  and  $x_j$
3. Compute MWST using Kruskal's algorithm. Pick a root, orient edges away from the root.

Steps 1-3 dominate the complexity – they all take  $O(n^2)$  time

23

## References

- Chow, C. K. and Liu, C. N.  
“Approximating Discrete Probability Distributions with Dependence Trees”.
- Meila-Predoviciu, M. Learning with Mixtures of Trees, PhD Thesis, MIT, 1999.

24